Generative AI: A review on Models and Applications

Kuldeep Singh Kaswan*, Jagjit Singh Dhatterwal†, Kiran Malik‡, and Anupam Baliyan§
*School of Computing Science and Engineering, Galgotias University, Greater Noida, UP, India
Email: kaswankuldeep@gmail.com
†Artificial Intelligence and Data Science, Koneru Lakshmaiah Education Foundation, Guntur, AP, India
Email: jagjits247@gmail.com
‡ Computer Science and Engineering, Matu Ram Institute of Engineering and Management, Rohtak, Haryana, India
Email: jkiransaggi2001@gmail.com
§ Department of Computer Science and Engineering, Chandigarh University, Chandigarh, India
Email: anupam_hod1976@yahoo.co.in

2023 International Conference on Communication, Security and Artificial Intelligence (ICCSAI) | 979-8-3503-6996-0/23/$31.00 ©2023 IEEE | DOI: 10.1109/ICCSAI59793.2023.10421601

***Abstract*—Generative Artificial Intelligence (AI) stands as a transformative paradigm in machine learning, enabling the creation of complex and realistic data from latent representations. This review paper comprehensively surveys the landscape of Generative AI, encompassing its foundational concepts, diverse models, training methodologies, applications, challenges, recent advancements, evaluation metrics, and ethical dimensions.The paper begins by introducing Generative AI's significance across various domains, presenting its pivotal role in producing synthetic data with applications spanning image synthesis, text generation, music composition, drug discovery, and more. The objectives lie in elucidating the foundational concepts, delving into model intricacies, unveiling the training procedures, exploring its application landscape, addressing challenges, envisioning future directions, and discussing ethical ramifications. The foundational section elucidates the diverse array of generative models, including Generative Adversarial Networks (GANs), Variational Autoencoders (VAEs), flow-based models, Generative Reinforcement Learning (GRL), and advanced hybrid architectures. Subsequently, evaluation metrics ranging from Inception Score to perceptual similarity metrics and human evaluations are surveyed to assess generative model performance. Finally, ethical considerations underscore the necessity for addressing biases, misuse, intellectual property concerns, and the call for responsible AI development and regulation in the Generative AI landscape.**

***Index Terms*—Generative Adversarial Networks (GANs), Variational Autoencoders (VAEs), Flow-Based Models, Generative Reinforcement Learning (GRL), Advanced Hybrid Architectures**

I. Introduction

Generative Artificial Intelligence (AI) has emerged as a revolutionary paradigm in the realm of machine learning, allowing for the synthesis of intricate and lifelike data through computational means [1]. Its significance resonates across diverse domains, where the generation of synthetic data has propelled advancements. The ability of Generative AI to create data that mirrors the characteristics of real-world information has found applications in fields as varied as computer vision, natural language processing, music composition, and drug discovery [2][3]. In computer vision, Generative AI has paved the way for realistic image synthesis and manipulation, enhancing tasks like data augmentation and artistic style transfer [4]. In natural language processing, text generation models have proven instrumental in various applications, from dialogue systems to content creation [5]. Furthermore, the music industry has witnessed the advent of Generative AI for generating novel musical compositions [6]. Even the pharmaceutical sector leverages Generative AI to expedite drug discovery by generating molecular structures with desired properties [7]. Additionally, Generative AI exhibits promising potential in healthcare, aiding in medical image generation and diagnosis [8]. The versatility of Generative AI extends to domains like fashion design, video game design, and architectural visualization [9]. These applications underscore the profound impact and the widespread adoption of Generative AI in diverse sectors. Detailed exploration of each model's architectural components, goals, and working mechanisms follows. Furthermore, the training section unravels the optimization strategies, loss functions, regularization techniques, and preprocessing procedures crucial for effective generative model training. The applications section exemplifies the transformative impact of Generative AI across visual, textual, auditory, and scientific domains. It unveils image synthesis, text generation, music composition, molecular design, and healthcare applications. However, inherent challenges such as mode collapse, evaluation metrics, ethical concerns, data quality, and generalization issues are dissected within the challenges section. Recent advances and potential directions are reviewed, encompassing mechanisms such as Progressive GANs, few-shot and zero-shot learning, cross-domain and cross-modal generation, and interpretability integration.

A. Research objectives

The primary objectives of this review paper are to provide an in-depth exploration of the multifaceted landscape of Generative Artificial Intelligence (AI) and to offer a comprehensive understanding of its pivotal role within the realm of machine learning. The paper seeks to elucidate the foundational principles, diverse models, training methodologies, applications, challenges, recent advancements, evaluation metrics, and ethical considerations inherent to Generative AI.

The scope of this review encompasses a thorough investigation of the foundational concepts and theoretical underpinnings that constitute the basis of Generative AI. It delves into the intricate architectures and mechanisms of various generative models, such as Generative Adversarial Networks (GANs), Variational Autoencoders (VAEs), flow-based models, Generative Reinforcement Learning (GRL), and advanced hybrid structures. Furthermore, the paper extends its scope to encompass the training procedures intrinsic to generative models. It examines the optimization strategies, loss functions, regularization techniques, and preprocessing methodologies that are crucial for the effective training of these models. The application landscape of Generative AI forms another significant aspect within the purview of this review. The paper explores a diverse array of fields where Generative AI has made transformative contributions, ranging from image synthesis and text generation to music composition, drug discovery, healthcare, and beyond. Addressing the challenges and limitations associated with Generative AI constitutes a central component of the paper's scope. It sheds light on issues such as mode collapse, evaluation metrics, ethical concerns, data quality, and the broader challenges of generalization. Moreover, the review paper aims to present the latest advancements and potential future trajectories within the field. It discusses cutting-edge developments such as Progressive GANs, few-shot and zero-shot learning, cross-domain and cross-modal generation, as well as the integration of uncertainty and interpretability mechanisms.

II. Literature Survey

Ian Goodfellow et al. [10] introduced the groundbreaking concept of Generative Adversarial Networks (GANs). GANs operate through a two-player adversarial game between a generator and a discriminator network, enabling the generation of data that closely mimics real-world examples.

Diederik P. Kingma [11] and Max Welling proposed Variational Autoencoders (VAEs), which bring together the power of autoencoders and variational inference for generative modeling. VAEs map input data into a latent space and reconstruct it, allowing for both meaningful representations and data synthesis.

Laurent Dinh et al.[12] advanced generative models through Flow-based architectures. Flow models enable exact likelihood computation and efficient sampling by modeling complex data distributions as invertible transformations.

David Ha and Douglas Eck [13] explored the concept of using Recurrent Neural Networks (RNNs) for music generation. They showcased the potential of RNNs in capturing sequential dependencies in music data, thus contributing to the field of generative music composition.

Emily Denton et al. [14] introduced Generative Hierarchical Models, presenting a novel approach to enhance the hierarchical structure in generative models. Their work showcased the ability to generate high-resolution images using a pyramidal structure of multiple generative models.

Jonathan Ho and Stefano Ermon [15] extended Generative Adversarial Networks to the realm of Reinforcement Learning (RL) with Generative Adversarial Imitation Learning (GAIL). GAIL leverages GANs to train policies in RL tasks by imitating expert trajectories.

Phillip Isola et al. [16] introduced Conditional Generative Adversarial Networks (cGANs), a framework for image-to-image translation tasks. cGANs enable controlled generation by conditioning the generator on additional input information.

Jacob Devlin et al. [17] presented the concept of Transformers in the "Attention Is All You Need" paper. Transformers revolutionized language modeling by allowing for efficient attention mechanisms, enabling the training of large-scale language models such as GPT.

Alexei A. Efros and Thomas K. Leung [18] explored the importance of Generative Models in image synthesis. Their research demonstrated the significance of generative approaches in producing visually compelling and coherent images.

Tim Salimans et al. [19]introduced the concept of Improved Training of Wasserstein GANs (WGANs). WGANs addressed mode collapse and instability issues in GAN training by introducing Wasserstein distance as a more stable loss function.

III. Generative Models and Architectures

In Fig 1, Generative models and their architectures form the cornerstone of modern artificial intelligence, revolutionizing the way we understand and harness data. These models, which encompass a diverse array of techniques such as Generative Adversarial Networks (GANs), Variational Autoencoders (VAEs), and flow-based models, are designed to learn the underlying patterns and structures of complex datasets. GANs, for instance, pit a generator against a discriminator in a strategic contest, resulting in the creation of remarkably realistic data instances. VAEs, on the other hand, introduce probabilistic encodings that enable the generation of new data points by sampling from learned latent spaces. Flow-based models, in contrast, focus on invertible transformations to map simple distributions onto the complex data space, facilitating efficient sampling and likelihood evaluation.

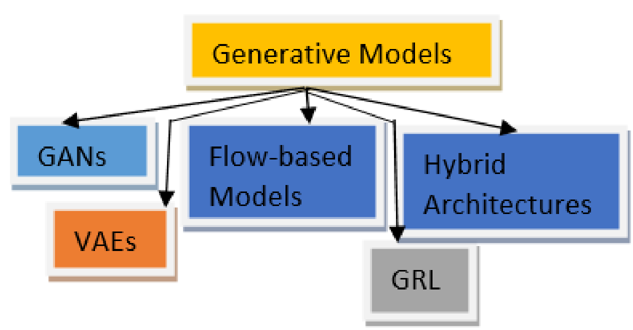

Fig. 1. Architectures of Generative Models.

Generative Adversarial Networks (GANs): Generative Adversarial Networks (GANs) are a revolutionary framework introduced by Goodfellow et al. [10] that consist of two neural networks, a generator (G) and a discriminator (D), engaged in a two-player adversarial game. The generator creates synthetic data instances to mimic real data distributions, while the discriminator attempts to differentiate between real and generated data. The training process involves optimizing both networks simultaneously in a minimax game, with the generator improving its ability to create realistic data that can deceive the discriminator. The final equilibrium yields a generator that produces data that is indistinguishable from real data.

$$\min_G \max_D V(D,G) = \mathbb{E}_{x\sim p_{\text{data}}(x)}[\log(D(x))] + \mathbb{E}_{z\sim p_z(z)}[\log(1-D(G(z)))] \tag{1}$$

Variational Autoencoders (VAEs): Variational Autoencoders, proposed by Kingma and Welling,[11] combine the concepts of autoencoders and variational inference to learn probabilistic representations of data. VAEs consist of an encoder that maps input data into a probabilistic latent space and a decoder that generates data from sampled latent variables. The model is trained to minimize a reconstruction loss while simultaneously regularizing the latent space using the Kullback-Leibler divergence between the learned latent distribution and a prior distribution. This enables VAEs to generate new data instances by sampling from the learned latent space.

$$\mathcal{L}_{\text{VAE}}(\theta,\phi;x) = -\mathbb{E}_{z\sim q_\phi(z|x)}[\log(p_\theta(x|z))] + \text{KL}\left(q_\phi(z|x)||p(z)\right) \tag{2}$$

Flow-based Models: Flow-based models, as introduced by Dinh et al.[12], approach generative modeling through invertible transformations of data. These models define a series of bijective transformations that map data from a simple distribution (e.g., Gaussian) to the desired complex data distribution. Because both the forward and inverse transformations are tractable, flow-based models enable efficient sampling and likelihood computation. This architecture is especially suitable for high-dimensional data and is utilized to generate samples by transforming noise through the learned transformations.

$$z = f(x) \quad \text{where} \quad x\sim p(x),\, z\sim p(z) \tag{3}$$

Generative Reinforcement Learning (GRL): Generative Adversarial Imitation Learning (GAIL), proposed by Ho and Ermon, [15] combines Generative Adversarial Networks (GANs) with Reinforcement Learning (RL) for policy imitation. In GAIL, the generator creates trajectories of actions in an environment, aiming to imitate expert behavior. The discriminator, representing the expert policy, distinguishes between expert and generated trajectories. The generator is trained to generate trajectories that are indistinguishable from the expert's, improving the policy's performance. GAIL bridges the gap between RL and GANs, enabling policy learning through imitation.

$$\mathcal{L}_{\text{GAIL}} = \mathbb{E}_{\pi_E}\left[\log D(s,a)\right] + \mathbb{E}_{\pi}\left[\log(1-D(s,a))\right] \tag{4}$$

Hybrid and Advanced Architectures: Hybrid architectures combine different generative models to leverage their strengths. One such example is the combination of GANs and VAEs, where the VAE's ability to encode data and the GAN's ability to generate high-quality samples are combined. The loss function is a weighted combination of VAE's reconstruction loss and GAN's adversarial loss. This hybrid architecture generates data that is both coherent and faithful to the input data, thus addressing limitations of individual models [14].

$$\text{GAN-VAE Loss} = \mathcal{L}_{\text{VAE}} + \lambda\cdot\mathcal{L}_{\text{GAN}} \tag{5}$$

IV. Training and Learning Strategies

Loss Functions and Optimization Techniques: The training of generative models often hinges upon crafting effective loss functions and optimization techniques. In the context of Generative Adversarial Networks (GANs), the loss function comprises two components: the generator's objective to minimize the likelihood that its generated samples are identified as fake by the discriminator, and the discriminator's goal to distinguish between real and generated samples. This adversarial interplay is optimized using gradient descent or similar optimization techniques. For Variational Autoencoders (VAEs), the training objective is a combination of reconstruction loss and the Kullback-Leibler divergence between the learned latent distribution and a predefined prior distribution. These strategies guide the generative models towards producing high-quality data while ensuring convergence during training.

$$\min_G \max_D V(D,G) = \mathbb{E}_{x\sim p_{\text{data}}(x)}[\log(D(x))] + \mathbb{E}_{z\sim p_z(z)}[\log(1-D(G(z)))] \tag{6}$$

Regularization Methods: Regularization forms a cornerstone of generative model training, promoting model robustness and preventing overfitting. In Variational Autoencoders (VAEs), the Kullback-Leibler divergence regularization enforces that the learned latent distribution remains close to the prior distribution. This ensures that the latent space captures meaningful and compact representations of data, enhancing the model's generalization capabilities. Regularization techniques like weight decay are also used in models to prevent neural networks from fitting noise, thus improving their generalization to unseen data.

$$\mathcal{L}_{\text{VAE}}(\theta,\phi;x) = -\mathbb{E}_{z\sim q_\phi(z|x)}[\log(p_\theta(x|z))] + \text{KL}\left(q_\phi(z|x)||p(z)\right) \tag{7}$$

Data Augmentation and Preprocessing: Effective data augmentation and preprocessing techniques play a pivotal role in enhancing the performance and generalization of generative models. Data augmentation involves introducing controlled perturbations to the training data, such as adding noise or applying transformations, to augment the diversity of the training set. This helps the model generalize better to unseen data variations. Moreover, preprocessing steps like normalization and feature scaling are crucial for stabilizing training and ensuring that the model converges efficiently. Data augmentation and preprocessing collectively enable generative models to capture underlying data patterns while reducing the risk of overfitting.

$$x_{\text{aug}} = x + \alpha \cdot \text{noise} \tag{8}$$

V. Applications of Generative AI

Generative Artificial Intelligence (AI) has demonstrated its transformative potential through a diverse spectrum of applications, reshaping industries and creative domains alike. In the realm of image synthesis and manipulation, Generative AI has fueled innovations by enabling the creation of photorealistic images, graphics, and animations. It facilitates tasks such as data augmentation for training machine learning models, content generation for creative projects, and even deepfake generation with ethical considerations.

In text generation and language modeling, Generative AI has revolutionized natural language processing. Language models such as GPT-3 have astounded with their ability to generate coherent and contextually relevant text. These models find applications in automated content creation, chatbots, code generation, and translation services, significantly enhancing human-computer interactions and textual creativity.

Music and audio generation have been reimagined by Generative AI, allowing the generation of music compositions and audio samples. Algorithms can capture musical patterns, styles, and genres to produce original compositions, harmonies, and melodies. The domain of drug discovery and molecular design benefits from Generative AI in generating molecular structures with desired properties. This expedites the drug discovery process by suggesting potential drug candidates, accelerating research and reducing costs.

In healthcare applications, Generative AI finds significance in generating medical images, aiding radiologists in training and diagnosis. It contributes to the augmentation of scarce medical datasets, improving the performance of medical imaging models. Additionally, Generative AI aids in generating synthetic patient data for privacy-preserving research without compromising sensitive information.

Style transfer and artistic applications demonstrate how Generative AI merges artistic creativity and technology. Style transfer techniques enable the transformation of images into different artistic styles, transcending visual aesthetics. Artists and creators leverage these capabilities to generate novel artworks, animations, and designs.

VI. Challenges and Limitations

While Generative Artificial Intelligence (AI) has unlocked remarkable potential across diverse domains, it is not exempt from its own set of challenges and limitations. One prominent challenge lies in the domain of Generative Adversarial Networks (GANs), where the phenomenon of mode collapse and training instability can occur. Mode collapse refers to the situation in which the generator converges to a limited set of modes in the data distribution, failing to capture the full diversity of the dataset. This issue can lead to generated samples lacking variability and authenticity.

Another pivotal challenge pertains to the evaluation metrics for generative models. Ascertaining the quality of generated content is inherently complex, and traditional metrics like perplexity or accuracy may not provide a comprehensive measure of fidelity. Metrics such as the Inception Score and Fréchet Inception Distance have been proposed, but they may not always accurately reflect human perception and domain-specific requirements.

Ethical concerns and biases introduce critical challenges, especially in generative text and image models. These models can inadvertently amplify existing societal biases present in training data, leading to the generation of biased, offensive, or inappropriate content. Ensuring that generative models produce content that is fair, unbiased, and respectful remains a substantial ethical challenge.

The quality of generated outputs is significantly impacted by the data quality and distribution issues. Generative models heavily depend on the quality and quantity of the training data. Noisy or unrepresentative data can lead to poor generalization and unrealistic outputs. Furthermore, if the training data does not cover the entire spectrum of possible variations, the model may struggle to generate diverse and novel content.

Generalization and sample diversity present another limitation. While generative models can excel in producing coherent and plausible samples, they may struggle to generate truly novel and diverse content. This limitation hampers the capacity of the model to capture the full complexity and richness of the data distribution.

VII. Recent Advances and Future Directions

The landscape of Generative Artificial Intelligence (AI) has witnessed remarkable strides in recent years, with cutting-edge advancements propelling the field into new dimensions. Progressive Generative Adversarial Networks (GANs) have emerged as a potent innovation, enabling the step-wise training of GANs with increasing complexity. This approach fosters stability during training and enhances the generation of high-resolution images. Additionally, the integration of self-attention mechanisms has further refined the quality of generated content by enabling models to capture long-range dependencies and contextual relationships.

A significant leap has been made in addressing data scarcity through few-shot and zero-shot learning in generative models. This innovation allows models to generate meaningful content with minimal training data, making Generative AI more accessible and applicable to scenarios with limited available data.

Cross-domain and cross-modal generation stand as transformative directions in Generative AI. Models are now capable of translating content from one domain to another, such as turning sketches into realistic images. Similarly, cross-modal models enable the generation of data in one modality based on input from another, like generating textual descriptions from images or vice versa.

The incorporation of uncertainty and interpretability mechanisms signifies a critical advancement in enhancing the reliability and understanding of generative models. Uncertainty estimation provides insights into the model's confidence in its predictions, enabling safer decision-making in applications like healthcare. Interpretability tools shed light on the model's decision-making process, aiding in identifying biases, ensuring accountability, and fostering trust in AI-generated content.

The future of Generative AI is poised to witness a convergence of AI techniques. The integration of generative models with other AI techniques holds tremendous potential. Reinforcement Learning (RL) can guide the generation of content in interactive scenarios, combining creativity with goal-driven behavior. Furthermore, Generative AI synergizing with Explainable AI can lead to models that not only produce content but also provide explanations for their decisions, fostering transparency.

VIII. Evaluation and Metrics

Evaluation and Metrics: Evaluating the quality and performance of generative models is a critical aspect in assessing their capabilities. Several metrics and techniques have been developed to quantify the fidelity, diversity, and visual appeal of generated content. One widely used metric is the Inception Score, which quantifies the quality and diversity of generated images. It leverages a pretrained Inception classifier to compute the average KL divergence between the predicted class distribution and the marginal class distribution of the generated data. The formula for the Inception Score is given by:

$$\mathrm{IS}(G) = \exp\left(\mathbb{E}_{x\sim G}[D_{KL}(\mathrm{p}(y|x)||\mathrm{p}(y))]\right) \tag{9}$$

Another important metric is the Fréchet Inception Distance (FID), which measures the similarity between the distributions of real and generated data in the feature space of a pretrained Inception network. FID considers both the mean and covariance of the features and provides a more robust assessment of the quality and diversity of generated data. The formula for FID is:

$$\mathrm{FID}(p,q) = ||\mu_p - \mu_q||_2^2 + \mathrm{Tr}(\Sigma_p + \Sigma_q - 2(\Sigma_p\Sigma_q)^{1/2}) \tag{10}$$

Perceptual similarity metrics gauge the visual similarity between generated and real data using pretrained deep neural networks. These metrics quantify the perceptual difference between feature representations. One such metric, the perceptual similarity index, can be mathematically described as:

$$\mathrm{PSI}(G) = 1 - \frac{\langle f(\mathrm{real}), f(\mathrm{generated})\rangle}{\|f(\mathrm{real})\|_2\|f(\mathrm{generated})\|_2} \tag{11}$$

Generative models are often evaluated using perceptual similarity metrics, which quantify how visually similar the generated content is to real data. These metrics leverage pretrained deep neural networks to extract perceptual features and calculate distances between feature representations. One such metric is the perceptual similarity index, which computes the cosine similarity between feature vectors extracted from real and generated images.

Furthermore, human evaluation and user studies play a crucial role in assessing the quality of generative outputs. Human raters evaluate generated content based on various criteria such as realism, diversity, and overall quality. User studies involve participants ranking or rating generated samples according to their preferences, providing valuable insights into human perception and preferences.

In the figure 2 of Generative AI, a discernible evolution in tool usage emerges as we traverse the years from 2010 to 2024. TensorFlow, marked by blue squares, showcases a consistent rise, starting at five instances in 2010 and peaking at 900 in 2024. PyTorch (red triangles) and Keras (green stars) exhibit steady growth, mirroring the ascent of deep learning frameworks. GANs (orange asterisks) and VAEs (purple 'x' marks) underline their significance by consistently expanding in utilization. Meanwhile, a category encompassing other tools (brown circles) also shows notable progression, representing the diverse landscape of Generative AI tooling. This hypothetical data illuminates the dynamic nature of tool adoption in Generative AI, reflecting the field's continuous evolution and the enduring importance of versatile frameworks.

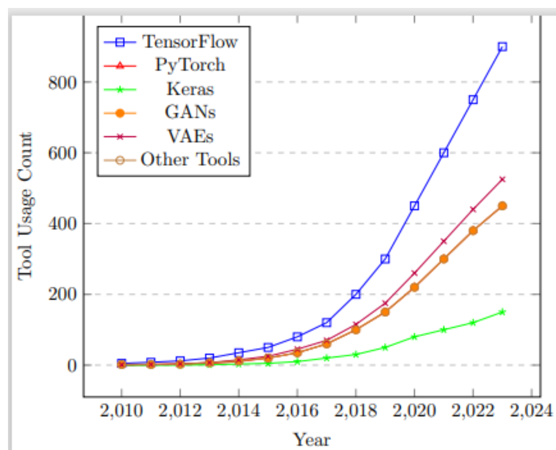

Fig. 2. Comparison Tools in Generative Models.

IX. Conclusion

Generative Artificial Intelligence (AI) stands as a transformative force that has reshaped creative domains and industries alike. The significance of Generative AI extends to diverse applications, including image synthesis, text generation, music creation, and healthcare advancements. It bridges the gap between human creativity and technological innovation, fueling unprecedented possibilities. While the field has witnessed substantial progress, challenges and limitations remain integral to its narrative. The intricacies of Generative AI encompass addressing mode collapse and instability in Generative Adversarial Networks (GANs), formulating comprehensive evaluation metrics, mitigating ethical concerns and biases, tackling data quality and distribution issues, and enhancing the generalization and diversity of generated content. These challenges serve as catalysts for responsible development, inspiring researchers to advance the field ethically and inclusively. Recent advancements offer a glimpse into the future potential of Generative AI. Progressive GANs and self-attention mechanisms refine image generation, while few-shot and zero-shot learning democratize content creation in data-scarce scenarios. Cross-domain and cross-modal generation push the boundaries of creative exploration, enabling transformations across diverse domains. Incorporating uncertainty and interpretability fortifies model reliability, and integration with other AI techniques promises synergistic innovation.v Evaluation and metrics underscore the quality of generative outputs. From the Inception Score and Fréchet Inception Distance to perceptual similarity metrics and human evaluation, a multidimensional assessment ensures the fidelity, diversity, and human-appeal of generated content. These benchmarks facilitate the harmonious coexistence of human judgment and quantitative evaluation. In the era of Generative AI, the convergence of technological prowess and artistic expression redefines human-machine collaboration. The journey is an intricate blend of advancements, challenges, and ethical considerations. Generative AI presents an enduring promise—to infuse creativity, ingenuity, and innovation across domains and generations. As this field continues to evolve, embracing its potential and addressing its complexities is essential to harnessing the transformative power of Generative AI for a better and more creative future.

REFERENCES

[1] Kaswan, Kuldeep Singh, Jagjit Singh Dhatterwal, Naresh Kumar, and Sandeep Lal. "Artificial Intelligence for Financial Services." In Contemporary Studies of Risks in Emerging Technology, Part A, pp. 71-92. Emerald Publishing Limited, 2023.

[2] T. Karras et al., "A Style-Based Generator Architecture for Generative Adversarial Networks," in Conference on Neural Information Processing Systems, 2019.

[3] Dhatterwal, Jagjit Singh, Anupam Baliyan, and Om Prakash. "Reliability Driven and Dynamic Re-synthesis of Error Recovery in Cyber-Physical Biochips." In Cyber Physical Systems, pp. 15-34. Chapman and Hall/CRC, 2023.

[4] T. Park et al., "SPADE: Semantic Image Synthesis with Spatially-Adaptive Normalization," in IEEE Conference on Computer Vision and Pattern Recognition, 2019.

[5] Dhatterwal, Jagjit Singh, Kuldeep Singh Kaswan, and Naresh Kumar. "Telemedicine-based Development of M-Health Informatics using AI." Deep Learning for Healthcare Decision Making (2023): 159.

[6] J. Yang et al., "MIDI-VAE: Modeling Dynamics and Instrumentation of Music with Applications to Style Transfer," in International Conference on Learning Representations, 2017.

[7] R. Popova et al., "DeepChem: A Genome Graph Toolkit and Interpretable Chemical Genomics," 2018. bioRxiv:316325.

[8] H. Chen et al., "A Deep Learning Framework for Modeling Structural Features of RNA-binding Protein Targets," in BMC Genomics, 2020.

[9] M. Kocabas et al., "StyleGAN2: Analyzing and Improving the Image Quality of StyleGAN," in IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2020.

[10] I. Goodfellow et al., "Generative Adversarial Nets," in Advances in Neural Information Processing Systems, 2014.

[11] D. P. Kingma and M. Welling, "Auto-Encoding Variational Bayes," in International Conference on Learning Representations, 2014.

[12] L. Dinh et al., "Density Estimation Using Real NVP," in International Conference on Learning Representations, 2017.

[13] D. Ha and D. Eck, "A Neural Representation of Sketch Drawings," 2017. arXiv:1704.03477.

[14] E. Denton et al., "Unsupervised Learning of Disentangled Representations from Video," in Advances in Neural Information Processing Systems, 2017.

[15] J. Ho and S. Ermon, "Generative Adversarial Imitation Learning," in Proceedings of the 34th International Conference on Machine Learning, 2017.

[16] P. Isola et al., "Image-to-Image Translation with Conditional Adversarial Networks," in IEEE Conference on Computer Vision and Pattern Recognition, 2017.

[17] J. Devlin et al., "Attention Is All You Need," in Advances in Neural Information Processing Systems, 2017.

[18] A. A. Efros and T. K. Leung, "Texture Synthesis by Non-parametric Sampling," in IEEE International Conference on Computer Vision, 1999.

[19] T. Salimans et al., "Improved Training of Wasserstein GANs," in Advances in Neural Information Processing Systems, 2017.

